

Revaltus Platform

How the system works, what it costs, where it can break, and who runs it

2 code repositories	9 pipeline stages	3 Claude models	\$10/mo per client site	20+ clients/mo target	1 developer (Hank)
----------------------------------	-----------------------------	---------------------------	-----------------------------------	---------------------------------	---------------------------------

Executive summary

Revaltus turns a CPA firm's existing website into a finished, freshly written new site through a single production platform plus a reusable site template. The platform audits a prospect's site, builds an AI “Master Firm Profile,” runs a short client onboarding chat, then generates a full content package and publishes it into the client's own website repository. A person reviews and approves at the key gates; the AI does the heavy drafting; a producer finishes and deploys the site.

- **Two systems, one flow.** A **platform** (audit + onboarding + content generation) and a per-client **website template** (the theme). They connect through a content-package zip.
- **AI does drafting, humans own accuracy.** Claude writes copy; people must review outlines and final copy for correctness and voice — the automated checks catch banned phrases, not wrong facts.
- **Client sites are static.** Generated sites are flat files on Vercel with no database, ~\$10/mo each, rebilled to clients.
- **Single-developer dependency.** Hank is the only person who can build or fix the platform and template today — the biggest operational risk.

Note: this supersedes the older [AI-Onboarding-Agent-Overview.pdf](#), which described a Basecamp-based intake-only tool. Basecamp has been removed; the audit, content-generation, and GitHub publishing capabilities are new.

1 · What the system is

A. Production platform — counting-five-onboarding

Next.js 16 / React 19 / TypeScript on Vercel, backed by Supabase (Postgres + Auth + private file storage). One app performs four jobs:

- **Site audits** — crawl up to 50 pages, Google PageSpeed, keyword checks, WHOIS, 9-category scoring, branded report with score-delta history.
- **Master Firm Profile (MBP)** — an AI intelligence brief that pre-fills 70–80% of intake data.
- **Client onboarding chat** — UUID-link conversation that confirms known data and fills gaps in 5–7 minutes.
- **Content pipeline** — six admin-gated phases that research, outline, write and package a full site, then seed the client's GitHub repo.

B. Client website template — counting-five-client-template

A second Next.js 16 app that is the actual *theme*. Every client gets their own git clone (GitHub [HankPantier/CountingFiveTemplate](#)).

- **25-block component library** + assembly system that turns generated .md files into rendered pages.

- **Auto-theming:** one unpack command ingests the content zip and generates a WCAG-corrected theme from the brand/design JSON.
- **Static output, no database.** Ships flat files; forms use Resend or an external endpoint; analytics via GA4/GTM; spam via Vercel BotID.
- **Not a shared dependency.** 20 clients = 20 repos = 20 Vercel projects. Template improvements must be pulled into each clone.

2 · Steps to onboard & ship one client

Skill level to run each stage: **[Low]** simple/operator · **[Moderate]** editor/producer · **[Senior dev]** developer.

- **1. Audit (internal, optional) [Low]**
Enter a prospect URL; background job scores the site and renders a report for the pitch.
- **2. Master Firm Profile [Low–Med]**
Generate the MBP (Sonnet 5) — the brief that pre-fills most intake data.
- **3. Session creation [Low]**
Upload MBP; app parses it to a structured schema and mints a unique client link.
- **4. Client chat (5–7 min) [Low]**
Client confirms known data, fills gaps, uploads assets. Auto-reminders after 3 idle days.
- **5. Admin review & approval [Low–Med]**
Review collected data and approve — the gate before any generation runs.
- **6. Content pipeline (6 phases) [Moderate]**
Lock palette → confirm sitemap → auto research → approve outlines → generate copy+SEO → assemble zip.
- **7. Publish the site [Mod–Tech]**
Seed client GitHub repo; producer unpacks, reviews theme, wires config/domain/env, deploys to Vercel.

*The two labor-heavy stages are **6 (content editing)** and **7 (site production & deploy)** — this is where per-client hours and error risk concentrate.*


3 · Where it lives

Layer	Service	Notes
Hosting	Vercel	Platform app + every client site. 300s Node functions; 2 cron jobs.
Database / Auth / Storage	Supabase	Postgres, admin auth, private asset bucket, RLS + capability gating. Platform only — client sites have no DB.
Source control	GitHub	Platform repo; client sites seeded from the template via a GitHub App.
AI — writing	Claude Sonnet 5	\$3 / \$15 per M tokens. All content generation.
AI — chat	Claude Sonnet 4.6	\$3 / \$15. Interactive intake (phases 3–4), audit/MBP/editor chats.
AI — cheap tasks	Claude Haiku 4.5	\$1 / \$5. Intake phases 1/2/5/6 + classification.
Other APIs	Serper, PageSpeed, Pexels, Resend, WHOIS	Keyword/SERP; page speed (free); stock photos (free); email; domain data.

All API/service accounts are agency-owned. The ~\$10/mo/client hosting is rebilled to clients.

4 · End-to-end workflow

From prospect audit to a live client website. The left column is the human-gated production pipeline; the right column runs automatically in parallel.



5 · Where it can break down

- **AI provider dependency.** An Anthropic outage, rate-limit, price change, or a model-ID rename (which silently prices at \$0) stalls content generation. Writing depends on Sonnet 5 specifically.
- **Content quality drift.** Automated checks catch banned phrases, not hallucinated facts, wrong services, or off-voice copy. Human editorial review is required every time.
- **Long-function timeouts.** A 20–30 page site can push the 300-second content-generation and packaging limits. Self-chaining + a stuck-job sweeper mitigate, but large sites are the risk zone.
- **The manual publish gap.** Stage 7 (repo → live site) is human-driven per client: theme review, image fixes, site config, env vars, domain, CSP. Highest-touch, highest-error step.
- **Two-repo drift.** The template evolves; already-shipped client sites don't get fixes unless someone pulls template updates into each clone.
- **Audit accuracy limits.** The crawler reads server HTML only, so JavaScript-heavy sites undercount. PageSpeed/Serper rate limits degrade results.
- **Secrets & config.** Service-role key, CRON_SECRET, GitHub App key must be correct and rotated. A bad env var breaks a whole stage. Supabase free tier won't hold 20+ clients of assets.
- **Key-person / bus factor.** Only Hank can build or fix the platform and template. Any absence halts development and breakage response.

6 · Staffing & support

Today Hank is the sole developer. Low-to-moderate roles are planned hires (TBD). Required coverage by function:

Function	Skill	Who / hire profile
Run audits	Low	Operator / account manager. No code. Marketing literacy to interpret for pitches.
MBP + onboarding sessions	Low–Med	Account manager / VA. Can run staff-proxy calls.
Content review & editing	Moderate	Editor with CPA-space literacy. The main recurring labor — accuracy + voice.
Site production & deploy	Mod–Tech	Front-end-capable web producer. unpack, config, domain/DNS, design polish.
Platform / template dev & fixes	Senior dev	Hank today. Next.js 16 / React 19 / TS / Supabase / Vercel / AI SDK. Bus-factor = 1.

Rough shape at 20+/mo: 1 developer (Hank) + 1–2 editors/producers + account/ops staff. Priority gap: a second developer or documented fallback to reduce single-person risk.

7 · Expenses & token projection

Figures are code-derived estimates. Token costs are modeled as ranges and should be validated against a real Anthropic bill.

Fixed monthly platform cost (agency)

Item	Est. / mo	Note
Vercel Pro (platform)	\$20+	Base + function compute on long jobs.
Supabase Pro	\$25	Free tier insufficient at volume.
Serper (keywords/SERP)	~\$50	Per-query; audit + content research.
Resend (email)	\$0–20	Free 3k/mo may suffice.
GitHub / PageSpeed / Pexels	\$0–4	Mostly free tiers.
Baseline subtotal	~\$100–150	Before tokens & client hosting.

Variable cost — Claude tokens (modeled)

Basis	Estimate	Driver
Onboarding chat / client	cents	Mostly Haiku, small exchanges.
Content generation / client	\$15–30	Sonnet 5, ~\$0.30–0.60/page × 20 pages + outlines/SEO/MBP. Thinking tokens swing 2–3×.
Anthropic at 20 clients/mo	~\$300–600/mo	Scales roughly linearly with clients & page count.

Client site hosting (rebilled)

- **~\$10/mo per live site**, static flat files on Vercel, **rebilled to clients** — revenue-neutral or margin, not a net cost.
- **Compounds with the install base:** hosting recurs per *live* site, not per new client. Adding 20 sites/mo adds ~\$200/mo of (rebilled) hosting each month — e.g. ~\$2,400/mo across ~240 sites after a year.

Bottom line: agency net variable cost is dominated by Claude tokens (~\$300–600/mo at 20 clients) on a ~\$100–150/mo fixed base. Client hosting is a pass-through. The largest real cost is staff labor, which appears on no API bill.

8 · Risks

AI / tech-stack

- Single-provider concentration on Anthropic; model deprecation or rename; price increases (Sonnet 5 intro pricing ends 2026-08-31).
- Hallucination / factual error requiring mandatory human review; a “silent \$0” pricing trap if a model ID drifts out of the pricing map.
- Bleeding-edge framework versions (Next 16, React 19) with a smaller support pool and faster breaking changes.

Operational / business

- Bus factor of 1 on development. No documented runbook for another dev to take over quickly.
- Two-repo drift: shipped client sites diverge from template fixes over time.

- **Scaling labor:** content editing and site production don't automate away — throughput is people-bound, not compute-bound.
- **Cost visibility:** token usage is logged but per-client billing rollups were still a recommended add — hard to attribute spend without it.

9 · Recommendations & assumptions

Priority actions

- **Reduce bus factor:** document a build/fix runbook and cross-train or contract a second developer.
- **Add per-client cost tracking** so token spend is attributable and billable.
- **Standardize stage 7** (site production) into a checklist/template to cut hours and errors as volume grows.
- **Template-update process** to propagate fixes into live client repos.
- **Validate token estimates** against a real Anthropic bill and set a spend alert.

Assumptions to confirm

- Token figures are modeled from code (budgets, models, page counts), not a live bill.
- Fixed-cost tiers (Vercel/Supabase/Serper) assume standard paid plans; confirm actual tiers.
- Client hosting rebilled at \$10/mo/site; all service accounts agency-owned (per your input).

Prepared for Revaltus internal ops planning · derived from the counting-five-onboarding and counting-five-client-template codebases.